

Artificial Intelligence

7. Logic Programming in Prolog

Prof. Bojana Dalbelo Bašić
Assoc. Prof. Jan Šnajder

University of Zagreb
Faculty of Electrical Engineering and Computing

Academic Year 2019/2020



Creative Commons Attribution-NonCommercial-NoDerivs 3.0

v1.10

Outline

- 1 Logic programming and Prolog
- 2 Inference on Horn clauses
- 3 Programming in Prolog
- 4 Non-declarative aspects of Prolog

Outline

- 1 Logic programming and Prolog
- 2 Inference on Horn clauses
- 3 Programming in Prolog
- 4 Non-declarative aspects of Prolog

Logic programming

- **Logic programming**: use of logic inference as a way of programming
- Main idea: define the problem in terms of logic formulae, then let the computer do the problem solving (program execution = inference)
- This is a typical **declarative programming** approach: express the logic of computation, don't bother with the control flow
- We focus on the description of the problem (**declarative**), rather than on how the program is executed (**procedural**)
- However, we still need some flow control mechanism, thus:

Algorithm = Logic + Control


- Different from automated theorem proving because:
 - ① explicit control flow is hard-wired into the program
 - ② not the full expressivity of FOL is supported

Refresher: Declarative programming

- Describe **what** is being computed instead of **how** (control flow)
- The programmer specifies a set of constraints that define the solution space, while finding the actual solution is left to the interpreter
- Hallmarks of declarative PL:
 - ▶ **explicit** rather than implicit state
 - ▶ **no side effects** or limited side effects
 - ▶ programming with **expressions**
- Two main flavors:
 - ▶ **functional PL**: expression is a function
 - ▶ **logic PL**: expression is a relation (represented by a predicate)
- **Upsides**: formally concise, high level of abstraction, lend themselves to formal analysis, less error prone
- **Downsides**: inefficiency, steep learning curve, no wide adoption

Prolog

- **Prolog** – “Programming in Logic”
- A declarative programming language
- “The offspring of a successful marriage between natural language processing and automated theorem-proving”
- Alan Colmerauer, Robert Kowalski, and Philippe Roussel in 1972



Colmerauer, A., & Roussel, P. (1996). [The Birth of Prolog](#). In History of programming languages (pp. 331–367)



SWI-Prolog's features

[HOME](#)[DOWNLOAD](#)[DOCUMENTATION](#)[TUTORIALS](#)[COMMUNITY](#)[USERS](#)[WIKI](#)

Overview

SWI-Prolog is a versatile implementation of the [Prolog](#) language. Although SWI-Prolog gained its popularity primarily in education, its development is mostly driven by the needs for **application development**. This is facilitated by a rich interface to other IT components by supporting many document types and (network) protocols as well as a comprehensive low-level interface to C that is the basis for high-level interfaces to C++, Java (bundled), C#, Python, etc (externally available). Data type extensions such as [dicts](#) and [strings](#) as well as full support for Unicode and unbounded integers simplify smooth exchange of data with other components.

SWI-Prolog aims at **scalability**. Its robust support for multi-threading exploits multi-core hardware efficiently and simplifies embedding in concurrent applications. Its *Just In Time Indexing* (JITI) provides transparent and efficient support for predicates with millions of clauses.

SWI-Prolog **unifies many extensions** of the core language that have been developed in the Prolog community such as *tabling*, *constraints*, *global variables*, *destructive assignment*, *delimited continuations* and *interactors*.

SWI-Prolog offers a variety of **development tools**, most of which may be combined at will. The native system provides an editor written in Prolog that is a close clone of Emacs. It provides *semantic* highlighting based on real time analysis of the code by the Prolog system itself. Complementary tools include a graphical debugger, profiler and cross-referencer. Alternatively, there is a mode for GNU-Emacs and, Eclipse plugin called [PDT](#) and a VSC [plugin](#), each of which may be combined with the native graphical tools. Finally, a *computational notebook* and web based IDE is provided by [SWISH](#). SWISH is a versatile tool that can be configured and extended to suit many different scenarios.

SWI-Prolog provides an add-on distribution and installation mechanism called **packs**. A *pack* is a directory with minimal organizational conventions and a *control* file that describes the origin, version, dependencies and automatic upgrade support. Packs

SWI Prolog

<https://swish.swi-prolog.org/>

The screenshot displays the SWISH web interface. The top menu bar includes 'SWISH', 'File', 'Edit', 'Examples', and 'Help'. Below the menu, there are tabs for 'houses_puzzle', 'examples', and 'kb'. The main editor area on the left contains the following Prolog code:

```
1 % Some simple test Prolog programs
2 % -----
3
4 % Knowledge bases
5
6 loves(vincent, mia).
7 loves(marcellus, mia).
8 loves(pumpkin, honey_bunny).
9 loves(honey_bunny, pumpkin).
10
11 jealous(X, Y) :-
12     loves(X, Z),
13     loves(Y, Z)
14
15
16 /** <examples>
17
18 ?- loves(X, mia).
19 ?- jealous(X, Y).
20
21 */
22
23
```

The right panel shows the execution trace for the query `trace, (jealous(X, Y)).`. The trace consists of several steps, each with a status (Call, Exit, Redo) and a message:

- Call: `jealous(_3972, _3976)`
- Call: `loves(_3972, _4402)`
- Exit: `loves(vincent, mia)`
- Call: `loves(_3976, mia)`
- Exit: `loves(vincent, mia)`
- Exit: `jealous(vincent, vincent)`
- X = Y, Y = vincent
- Redo: `loves(_3976, mia)`
- Exit: `loves(marcellus, mia)`
- Exit: `jealous(vincent, marcellus)`
- X = vincent,
- Y = marcellus
- Exit: `loves(marcellus, mia)`
- Call: `loves(_3976, mia)`
- Exit: `loves(vincent, mia)`
- Exit: `jealous(marcellus, vincent)`
- X = marcellus,
- Y = vincent

At the bottom of the trace panel, there are buttons for 'Next', '10', '100', '1,000', and 'Stop'. Below the trace panel, the query `?- trace, (jealous(X, Y)).` is entered in the input field.

Outline

- 1 Logic programming and Prolog
- 2 Inference on Horn clauses**
- 3 Programming in Prolog
- 4 Non-declarative aspects of Prolog

Horn clauses (1)

- Prolog uses a subset of FOL called **Horn clause logic**

Horn clause

A **Horn clause** is a clause (a disjunction of literals) with at most one positive literal:

$$\neg P_1 \vee \neg P_2 \vee \dots \vee \neg P_n \vee Q$$

or, equivalently:

$$(P_1 \wedge P_2 \wedge \dots \wedge P_n) \rightarrow Q$$

The negative literals constitute the **body** of the clause, while the positive literal constitutes its **head**.

Definite clause

A **definite clause** is a Horn clause with exactly one positive literal.

- A Horn clause can be propositional or first-order.

Horn clauses (2)

- A Prolog program is made up of a sequence of definite clauses
- Each definite clause defines a **rule of inference** or a **fact**
- A fact is simply a definite clause of the form $\text{True} \rightarrow Q \equiv Q$
- Inference rules and facts are an intuitive way of formalizing human knowledge
- A Horn clause with only negative literals is called a **goal clause**
- Given the logic program as input, the aim is to prove the goal clause using **refutation resolution**

Horn clauses (3)

- Horn clauses are of **restricted expressivity**: one can not transform every FOL formula into an equivalent Horn clause
 - ▶ E.g., $\neg P \rightarrow Q$ or $P \rightarrow (Q \vee R)$
- However, in practice, this turns out not be a severe limitation
- The upside of limiting ourselves to Horn clauses is that we can implement efficient reasoning procedures using **resolution** with either **forward chaining** or **backward chaining**
- Forward/backward chaining over Horn clauses is **complete**
- Propositional Horn clauses: time complexity of inference is **linear** in the number of clauses
- First-order Horn clauses: inference is still undecidable, but generally more efficient than in unrestricted FOL

Backward chaining

- Starting with a knowledge base of facts and rules (the logic program) Γ and the negated goal $\neg P$, we aim to derive a NIL clause

$\neg P$ resolves with $C_1 \in \Gamma$ and generates new negated goal $\neg P_2$

$\neg P_2$ resolves with $C_2 \in \Gamma$ and generates new negated goal $\neg P_3$

\vdots

$\neg P_k$ resolves with $C_k \in \Gamma$ and generates NIL

- This process can be viewed as a **state space search** where:
 - ▶ each state is the negated current goal
 - ▶ initial state: $\neg P$
 - ▶ goal state: the NIL clause
 - ▶ operator: resolving $\neg P_i$ with a clause C_j from Γ
- NB:** Horn clauses are closed under resolution: the resolvent of two Horn clauses is itself a Horn clause

Backward chaining – Example 1

- Logic program:

$$\begin{array}{lll} (1) & A & \equiv A \\ (2) & B & \equiv B \\ (3) & (A \wedge B) \rightarrow C & \equiv \neg A \vee \neg B \vee C \\ (4) & (C \vee D) \rightarrow E & \equiv \neg C \vee E \\ (5) & & \neg D \vee E \end{array}$$

- Goal: clause E
- Initial state: $\neg E$
- Step 1: Resolving goal $\neg E$ and clause (4), new goal is $\neg C$
- Step 2: Resolving goal $\neg C$ and clause (3), new goal is $\neg A \vee \neg B$
- Step 3: Resolving goal $\neg A \vee \neg B$ and clause (1), new goal is $\neg B$
- Step 4: Resolving goal $\neg B$ and clause (2), deriving NIL

Backward chaining – Example 2

- Logic program:

$$\begin{array}{lll} (1) & A & \equiv A \\ (2) & B & \equiv B \\ (3) & (A \wedge B) \rightarrow D & \equiv \neg A \vee \neg B \vee D \\ (4) & (C \vee D) \rightarrow E & \equiv \neg C \vee E \\ (5) & & \neg D \vee E \end{array}$$

- Goal: clause E
- Initial state: $\neg E$
- Step 1: Resolving goal $\neg E$ and clause (4), new goal is $\neg C$
- Step 2: **Backtracking to the most recent choice point**
- Step 3: Resolving goal $\neg E$ and clause (5), new goal is $\neg D$
- Step 4: Resolving goal $\neg D$ and clause (3), new goal is $\neg A \vee \neg B$
- Step 5: Resolving goal $\neg A \vee \neg B$ and clause (1), new goal is $\neg B$
- Step 6: Resolving goal $\neg B$ and clause (2), deriving NIL

Backward chaining – algorithm

- Non-deterministic algorithm for resolution over Horn clauses:

Backward chaining

```
function BackwardChaining( $P, \Gamma$ )  
  if  $P = \text{NIL}$  then return true  
   $L \leftarrow \text{SelectLiteral}(P)$   
   $C \leftarrow \text{SelectResolvingClause}(L, \Gamma)$   
  if  $C = \text{fail}$  then return false  
   $P' \leftarrow \text{resolve}(P, C)$   
  BackwardChaining( $P', \Gamma$ )
```

- SelectLiteral selects one literal from the clause P (a negated literal from the negated goal)
- SelectResolvingClause selects the clause from Γ whose head (which is the only positive literal of a Horn clause) resolves with L
- There can be many such clauses, thus the search branches here

Backward chaining in Prolog

- The clauses in Γ are ordered by the programmer (typically: more specific clauses come first, followed by more general clauses)
- The negative literals in each clause (i.e., the order of atoms in the antecedent) are also ordered by the programmer
- The state space search is carried out in **depth-first order**: when the negated goal P is resolved with clause C , the negative literals in C are placed **in the original order at the beginning** of P
- The `SelectLiteral` selects the **first** literal in P . Thus, P is implemented as a **stack** (LIFO)
- This proof strategy is known as **SLD (Selective Linear Definite clause) resolution**

Outline

- 1 Logic programming and Prolog
- 2 Inference on Horn clauses
- 3 Programming in Prolog
- 4 Non-declarative aspects of Prolog

Prolog facts and rules

$$\forall x (\text{HUMAN}(x) \rightarrow \text{MORTAL}(x)) \wedge \text{HUMAN}(\textit{Socrates})$$

```
mortal(X) :- human(X).    % rule
human(socrates).          % fact
```

- Variables are uppercased, predicate symbols are lowercased
- Implications are in the form consequent :- antecedent
- Variables are implicitly universally quantified
- Every line ends with a full stop

Prolog queries

```
?- human(socrates).    % Is Socrates human?  
true  
?- human(doughnut).    % Is a doughnut human?  
false  
?- mortal(socrates).   % Inference: is Socrates mortal?  
true  
?- mortal(X).           % Who is mortal?  
X = socrates  
true
```

Adding conditions to rules

$$\forall x((\text{MAMMAL}(x) \wedge \text{SPEAKS}(x)) \rightarrow \text{HUMAN}(x))$$

```
human(X) :- mammal(X), speaks(X). % comma denotes "and"
```

$$\forall x((\text{MAMMAL}(x) \wedge \text{SPEAKS}(x) \wedge \text{PAYS_TAXES}(x)) \rightarrow \text{HUMAN}(x))$$

```
human(X) :-  
    mammal(X),  
    speaks(X),  
    pays_taxes(X).
```

Adding disjunctions to rules

$$\forall x ((\text{HUMAN}(x) \vee \text{ALIVE}(x)) \rightarrow \text{MORTAL}(x))$$

- The disjunction in the rule condition can be written in two ways
- Either the rule is factored into separate clauses:

```
mortal(X) :- human(X).    % first clause  
mortal(X) :- alive(X).    % second clause
```

- Or a disjunction is introduced in the rule body:

```
mortal(X) :-  
    human(X); alive(X).    % semicolon denotes "or"
```

n -ary predicates

- Binary predicates model binary relations:

```
teacher(socrates, plato).    % Socrates is Plato's teacher
teacher(cratylus, plato).
teacher(plato, aristotle).
```

- A rule for disciple relation as the inverse of teacher relation:

```
disciple(X, Y) :- teacher(Y, X).
```

- A rule for defining that someone is taught by a teacher:

```
taught(X) :- disciple(X, Y).
```

- As we don't care about the value of Y , we can also write:

```
taught(X) :- disciple(X, _).
```

Query examples

```
?- teacher(X, plato).    % Who is Plato's teacher?
```

```
X = socrates
```

```
X = cratylus
```

```
true
```

```
?- teacher(socrates, Y), teacher(cratylus, Y).
```

```
    % Whom do they both teach?
```

```
Y = plato
```

```
true
```

```
?- taught(X).    % Who is being taught?
```

```
X = plato
```

```
X = plato
```

```
X = aristotle
```

```
true
```

```
?- disciple(aristotle, socrates).
```

```
false
```


Recursively defined predicates

- $\text{DISCIPLE}(x, y)$ captures only the direct relation
- How can we capture transitive relations?
- E.g., $\text{FOLLOWER}(x, y)$, iff x is an either direct or indirect follower of philosopher y :
 - ▶ Base case: x is an disciple of y
 - ▶ Recursive cause: x is an disciple of z , who in turn is the follower of y

```
follower(X, Y) :-    % base clause
    disciple(X, Y).
follower(X, Y) :-    % recursive clause
    disciple(X, Z),
    follower(Z, Y).
```

```
?- follower(aristotle, socrates).
true
```

Prolog search tree

- Applying the SLD resolution on the logic program and the negated goal generates a search tree which corresponds to a **proof tree**
- Each node is a **stack of negative literals** to be resolved
- The goal is to derive NIL, i.e., empty the stack
- Prolog attempts to resolve the top literal from the stack against the **head literal** of every clause from the program (remember: a literal from the stack is negative, while head literals from clauses are positive)
- Such literals are potentially **complementary unifiable**
- The clauses are searched from **top to bottom** of the program (hence the order of clauses is important)

Prolog search tree

- If the top literal L is complementary unifiable using MGU θ with the head of some clause C , then L gets popped from the stack, the body of C is pushed onto the stack, substitution θ is applied to all literals on the stack, and the search continues
- If no clause from the program is complementary unifiable with the L , the search **backtracks** to the last choice point \Rightarrow **depth-first search**
- If the stack is emptied, this means NIL is derived, hence return **true**
- If the search is complete and the stack is not emptied, return **false**

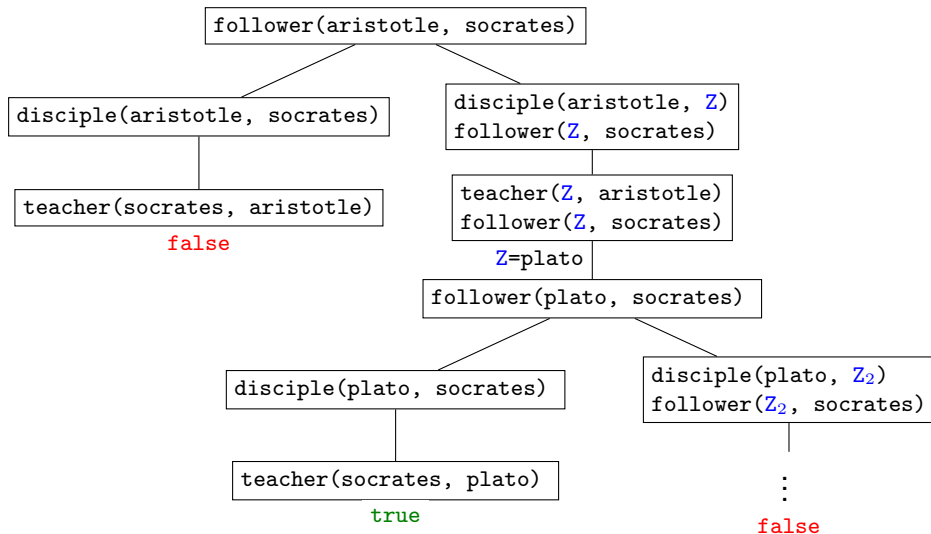
Prolog program – example

```
% Rules
follower(X, Y) :-
    disciple(X, Y).
follower(X, Y) :-
    disciple(X, Z),
    follower(Z, Y).

disciple(X, Y) :-
    teacher(Y, X).

% Facts
teacher(socrates, plato).
teacher(cratylus, plato).
teacher(plato, aristotle).
```

Prolog search tree – example



Prolog search tree – example – remarks

- Each node = downward-growing stack of negative literals
- The search branches on resolving the `follower` literal, because there are two rules for the `follower` predicate
- On request, after proving the goal, Prolog can be made to continue to search for alternative resolutions (which fails in this case)
- Do not confuse the stack of negative literals (constituting the state in the search space) with the DFS stack (used to implement the DFS; not shown on previous slide)

Prolog execution trace

```
[trace] ?- follower(aristotle, socrates).  
Call: (7) follower(aristotle, socrates) ? creep  
Call: (8) disciple(aristotle, socrates) ? creep  
Call: (9) teacher(socrates, aristotle) ? creep  
Fail: (9) teacher(socrates, aristotle) ? creep  
Fail: (8) disciple(aristotle, socrates) ? creep  
Redo: (7) follower(aristotle, socrates) ? creep  
Call: (8) disciple(aristotle, _G5025) ? creep  
Call: (9) teacher(_G5024, aristotle) ? creep  
Exit: (9) teacher(plato, aristotle) ? creep  
Exit: (8) disciple(aristotle, plato) ? creep  
Call: (8) disciple(plato, socrates) ? creep  
Call: (9) teacher(socrates, plato) ? creep  
Exit: (9) teacher(socrates, plato) ? creep  
Exit: (8) disciple(plato, socrates) ? creep  
Exit: (7) follower(aristotle, socrates) ? creep  
true.
```

Outline

- 1 Logic programming and Prolog
- 2 Inference on Horn clauses
- 3 Programming in Prolog
- 4 Non-declarative aspects of Prolog

Order of atoms/clauses

```
follower(X, Y) :-    % base clause
    disciple(X, Y).
follower(X, Y) :-    % recursive clause
    disciple(X, Z),
    follower(Z, Y).
```

- Formally, the order of clauses and their atoms should be arbitrary, due to commutativity of ' \wedge ' and ' \vee ':

$$\begin{aligned} & (\neg D(x, y) \vee F(x, y)) \wedge (\neg D(x, z) \vee \neg F(z, y) \vee F(x, y)) \\ \equiv & (\neg D(x, y) \vee F(x, y)) \wedge (\neg F(z, y) \vee \neg D(x, z) \vee F(x, y)) \\ \equiv & (\neg D(x, z) \vee \neg F(z, y) \vee F(x, y)) \wedge (\neg D(x, y) \vee F(x, y)) \\ \equiv & (\neg F(z, y) \vee \neg D(x, z) \vee F(x, y)) \wedge (\neg D(x, y) \vee F(x, y)) \end{aligned}$$

Order of atoms/clauses

- However, because Prolog uses SLD, commutativity doesn't hold and the order of clauses/atoms becomes important

1:

```
follower(X, Y) :-  
    disciple(X, Y).  
follower(X, Y) :-  
    disciple(X, Z),  
    follower(Z, Y).
```

2:

```
follower(X, Y) :-  
    disciple(X, Z),  
    follower(Z, Y).  
follower(X, Y) :-  
    disciple(X, Y).
```

3:

```
follower(X, Y) :-  
    disciple(X, Y).  
follower(X, Y) :-  
    follower(Z, Y),  
    disciple(X, Z)..
```

4:

```
follower(X, Y) :-  
    follower(Z, Y),  
    disciple(X, Z).  
follower(X, Y) :-  
    disciple(X, Y).
```

⇒ **declarative meaning** deviates from **procedural meaning**!

Negation

- A Horn clause does not allow for negated atoms in the antecedent, e.g.:

$$(Q(x) \wedge \neg P(x)) \rightarrow R(x) \equiv \neg Q(x) \underbrace{\vee P(x) \vee R(x)}_{\text{two pos. literals!}}$$

- Logic programming without negation would be far too restrictive
- Prolog introduces the `not` operator, which can be used in the body of a rule:

```
R(X) :- Q(X), not(P(X)).
```

- The semantics of `not` is different from the one in logic:

Negation as failure – NAF

The literal `not(P(x))` is true if `P(X)` cannot be derived, otherwise it is false.

- In logic: `not(P(x))` is true iff `P(x)` is false

Negation – example

```
human(X) :-  
    speaks(X),  
    not(has_feathers(X)).  
  
speaks(socrates).  
speaks(polynesia).  
has_feathers(polynesia).
```

```
?- human(polynesia).  
false  
?- human(socrates).  
true  
?- not(human(polynesia)).  
true
```

Closed world assumption

- NAF: if $P(x)$ can't be derived, then $\text{not}(P(x))$ true
- Standard semantics: if $\text{not}(P(x))$ is true, then $P(x)$ is false
- Taken together: if $P(x)$ can't be derived, then $P(x)$ is false
- In other words, **all things that can't be derived are false**

Closed world assumption – CWA

Everything that cannot be derived (facts that are not in the knowledge base and that cannot be derived from the knowledge base) is false.

- We don't allow for facts to be unknown (neither true nor false)
- CWA causes Prolog to deviate from standard semantics. E.g., in logic:

$$P, (P \wedge \neg Q) \rightarrow R \not\models R$$

but in Prolog:

$$P, (P \wedge \neg Q) \rightarrow R \vdash R$$

Wrap-up

- **Logic programming** is a kind of declarative programming, and **Prolog** is a logic programming language
- **Horn clause logic** is a subset of FOL which allows for efficient inference using resolution
- **Horn clauses** are clauses with at most one positive literal, while **definite clauses** are Horn clauses with exactly one positive literal
- A Prolog program is a sequence of definite first-order clauses, which correspond to **facts** and **rules**
- The program is executed by applying **refutation resolution with backward chaining** (SLD resolution)
- Non-commutativity of disjunction/conjunction and **negation as failure** are the non-declarative aspects of Prolog



Next topic: Expert systems